

Linear Regression Predictive Modeling on Bike Sharing Data:

Problem Statement

A bike-sharing system is a service in which bikes are made available for shared use to individuals on a short term basis for a price or free. Many bike share systems allow people to borrow a bike from a "dock" which is usually computer-controlled wherein the user enters the payment information, and the system unlocks it. This bike can then be returned to another dock belonging to the same system.

A US bike-sharing provider BoomBikes has recently suffered considerable dips in their revenues due to the ongoing Corona pandemic. The company is finding it very difficult to sustain in the current market scenario. So, it has decided to come up with a mindful business plan to be able to accelerate its revenue as soon as the ongoing lockdown comes to an end, and the economy restores to a healthy state.

In such an attempt, BoomBikes aspires to understand the demand for shared bikes among the people after this ongoing quarantine situation ends across the nation due to Covid-19. They have planned this to prepare themselves to cater to the people's needs once the situation gets better all around and stand out from other service providers and make huge profits.

They have contracted a consulting company to understand the factors on which the demand for these shared bikes depends. Specifically, they want to understand the factors affecting the demand for these shared bikes in the American market. The company wants to know:

Which variables are significant in predicting the demand for shared bikes. How well those variables describe the bike demands Based on various meteorological surveys and people's styles, the service provider firm has gathered a large dataset on daily bike demands across the American market based on some factors.

Business Goal:

You are required to model the demand for shared bikes with the available independent variables. It will be used by the management to understand how exactly the demands vary with different features. They can accordingly manipulate the business strategy to meet the demand levels and meet the customer's expectations. Further, the model will be a good way for management to understand the demand dynamics of a new market.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import RFE

import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

Reading The Data

```
In [3]: bike = pd.read_csv('day.csv')
#reading the data
bike
```

Out[3]:

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	registered	c	
0	1	01-01-2018	1	0	1	0	1	1	1	2	14.110847	18.18125	80.5833	10.749882	331	654	98
1	2	02-01-2018	1	0	1	0	2	1	2	14.902598	17.68695	69.6087	16.652113	131	670	80	
2	3	03-01-2018	1	0	1	0	3	1	1	8.050924	9.47025	43.7273	16.636703	120	1229	134	
3	4	04-01-2018	1	0	1	0	4	1	1	8.200000	10.60610	59.0435	10.739832	108	1454	150	
4	5	05-01-2018	1	0	1	0	5	1	1	9.305237	11.46350	43.6957	12.522300	82	1518	160	
...	
725	726	27-12-2019	1	1	12	0	5	1	2	10.420847	11.33210	65.2917	23.458911	247	1867	21	
726	727	28-12-2019	1	1	12	0	6	0	2	10.386653	12.75230	59.0000	10.416557	644	2451	30	
727	728	29-12-2019	1	1	12	0	0	0	2	10.386653	12.12000	75.2917	8.333661	159	1182	134	
728	729	30-12-2019	1	1	12	0	1	1	1	10.489153	11.58500	48.3333	23.500518	364	1432	179	
729	730	31-12-2019	1	1	12	0	2	1	2	8.849153	11.17435	57.7500	10.374682	439	2290	27	

730 rows × 16 columns



The Dataset has 730 rows and 16 columns

In [4]: bike.shape #size of the data

Out[4]: (730, 16)

In [5]: `bike.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   instant     730 non-null    int64  
 1   dteday      730 non-null    object  
 2   season      730 non-null    int64  
 3   yr          730 non-null    int64  
 4   mnth        730 non-null    int64  
 5   holiday     730 non-null    int64  
 6   weekday     730 non-null    int64  
 7   workingday  730 non-null    int64  
 8   weathersit  730 non-null    int64  
 9   temp         730 non-null    float64 
 10  atemp        730 non-null    float64 
 11  hum          730 non-null    float64 
 12  windspeed   730 non-null    float64 
 13  casual       730 non-null    int64  
 14  registered   730 non-null    int64  
 15  cnt          730 non-null    int64  
dtypes: float64(4), int64(11), object(1)
memory usage: 91.4+ KB
```

In [6]: `bike.describe()`

Out[6]:

	instant	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspe
count	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000
mean	365.500000	2.498630	0.500000	6.526027	0.028767	2.995890	0.690411	1.394521	20.319259	23.726322	62.765175	12.7636
std	210.877136	1.110184	0.500343	3.450215	0.167266	2.000339	0.462641	0.544807	7.506729	8.150308	14.237589	5.1958
min	1.000000	1.000000	0.000000	1.000000	0.000000	0.000000	0.000000	1.000000	2.424346	3.953480	0.000000	1.5002
25%	183.250000	2.000000	0.000000	4.000000	0.000000	1.000000	0.000000	1.000000	13.811885	16.889713	52.000000	9.0416
50%	365.500000	3.000000	0.500000	7.000000	0.000000	3.000000	1.000000	1.000000	20.465826	24.368225	62.625000	12.1253
75%	547.750000	3.000000	1.000000	10.000000	0.000000	5.000000	1.000000	2.000000	26.880615	30.445775	72.989575	15.6255
max	730.000000	4.000000	1.000000	12.000000	1.000000	6.000000	1.000000	3.000000	35.328347	42.044800	97.250000	34.0000



In [7]: bike.size

Out[7]: 11680

Data Cleaning & droping columns which are not relevant for analysis

In [8]: bike.isnull().sum()

```
Out[8]: instant      0
         dteday       0
         season       0
         yr           0
         mnth        0
         holiday      0
         weekday      0
         workingday   0
         weathersit    0
         temp          0
         atemp         0
         hum           0
         windspeed     0
         casual        0
         registered    0
         cnt           0
         dtype: int64
```

```
In [9]: bike.drop(['instant'],axis=1,inplace=True)
```

```
In [10]: bike.drop(['casual'],axis=1,inplace=True)
```

Inspecting data after dropping

```
In [11]: bike.head()
```

```
Out[11]:   dteday  season  yr  mnth  holiday  weekday  workingday  weathersit  temp  atemp  hum  windspeed  registered  cnt
0  01-01-2018      1    0     1      0       1        1         1        2  14.110847  18.18125  80.5833  10.749882      654    985
1  02-01-2018      1    0     1      0       2        1         1        2  14.902598  17.68695  69.6087  16.652113      670    801
2  03-01-2018      1    0     1      0       3        1         1        1   8.050924  9.47025  43.7273  16.636703     1229   1349
3  04-01-2018      1    0     1      0       4        1         1        1   8.200000  10.60610  59.0435  10.739832     1454   1562
4  05-01-2018      1    0     1      0       5        1         1        1   9.305237  11.46350  43.6957  12.522300     1518   1600
```

```
In [12]: bike.drop(['dteday'],axis=1,inplace=True)
```

In [13]: `bike.head()`

```
Out[13]:   season  yr  mnth  holiday  weekday  workingday  weathersit  temp  atemp  hum  windspeed  registered  cnt
0       1  0       1       0       1           1            2  14.110847  18.18125  80.5833  10.749882      654     985
1       1  0       1       0       2           1            2  14.902598  17.68695  69.6087  16.652113      670     801
2       1  0       1       0       3           1            1  8.050924   9.47025  43.7273  16.636703     1229    1349
3       1  0       1       0       4           1            1  8.200000  10.60610  59.0435  10.739832     1454    1562
4       1  0       1       0       5           1            1  9.305237  11.46350  43.6957  12.522300     1518    1600
```

In [14]: `bike.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 13 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   season      730 non-null    int64  
 1   yr          730 non-null    int64  
 2   mnth        730 non-null    int64  
 3   holiday     730 non-null    int64  
 4   weekday     730 non-null    int64  
 5   workingday  730 non-null    int64  
 6   weathersit  730 non-null    int64  
 7   temp         730 non-null    float64 
 8   atemp        730 non-null    float64 
 9   hum          730 non-null    float64 
 10  windspeed   730 non-null    float64 
 11  registered  730 non-null    int64  
 12  cnt         730 non-null    int64  
dtypes: float64(4), int64(9)
memory usage: 74.3 KB
```

In [15]: `bike.columns`

```
Out[15]: Index(['season', 'yr', 'mnth', 'holiday', 'weekday', 'workingday',
       'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'registered', 'cnt'],
       dtype='object')
```

```
In [16]: bike.corr()
```

	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	re
season	1.000000e+00	-3.279074e-16	8.310321e-01	-0.010868	0.004012	0.017868	0.021306	0.333361	0.342014	0.208220	-0.229607	(
yr	-3.279074e-16	1.000000e+00	-5.162656e-16	0.008195	0.000685	-0.011852	-0.050322	0.048789	0.047215	-0.112547	-0.011624)
mnth	8.310321e-01	-5.162656e-16	1.000000e+00	0.018905	-0.004059	-0.010414	0.045613	0.219083	0.226430	0.224937	-0.208013	(
holiday	-1.086804e-02	8.195345e-03	1.890483e-02	1.000000	-0.020145	-0.257009	-0.034395	-0.028764	-0.032703	-0.015662	0.006257	-()
weekday	4.012478e-03	6.852851e-04	-4.059002e-03	-0.020145	1.000000	0.001588	0.034216	0.044876	0.037964	0.008780	-0.017230	(
workingday	1.786841e-02	-1.185197e-02	-1.041372e-02	-0.257009	0.001588	1.000000	-0.026332	0.002044	0.010657	0.053770	-0.002453	(
weathersit	2.130636e-02	-5.032247e-02	4.561335e-02	-0.034395	0.034216	-0.026332	1.000000	-0.119503	-0.120559	0.590277	0.039769	-()
temp	3.333607e-01	4.878919e-02	2.190833e-01	-0.028764	0.044876	0.002044	-0.119503	1.000000	0.991696	0.128565	-0.158186	(
atemp	3.420139e-01	4.721519e-02	2.264302e-01	-0.032703	0.037964	0.010657	-0.120559	0.991696	1.000000	0.141512	-0.183876	(
hum	2.082196e-01	-1.125471e-01	2.249368e-01	-0.015662	0.008780	0.053770	0.590277	0.128565	0.141512	1.000000	-0.248506	-()
windspeed	-2.296069e-01	-1.162435e-02	-2.080131e-01	0.006257	-0.017230	-0.002453	0.039769	-0.158186	-0.183876	-0.248506	1.000000	-()
registered	4.103102e-01	5.969106e-01	2.919516e-01	-0.109142	0.199326	0.005466	-0.259025	0.539436	0.543678	-0.089212	-0.217914	.
cnt	4.045838e-01	5.697285e-01	2.781909e-01	-0.068764	0.036183	-0.027640	-0.295929	0.627044	0.630685	-0.098543	-0.235132	(

Handling outliers

```
In [17]: bike.columns
```

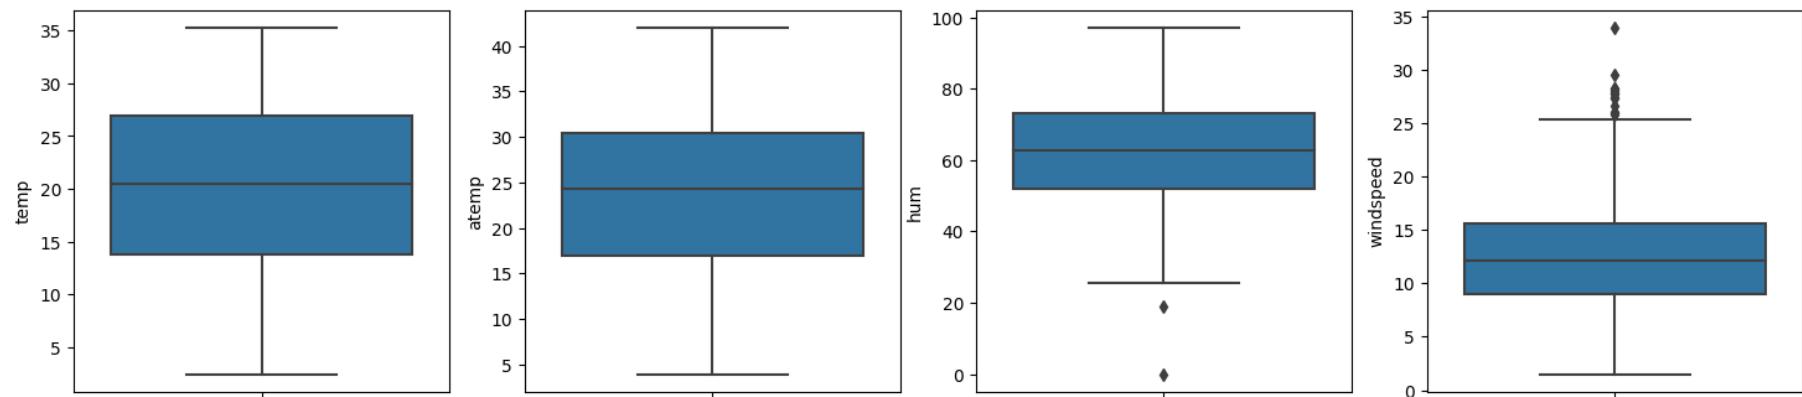
```
Out[17]: Index(['season', 'yr', 'mnth', 'holiday', 'weekday', 'workingday',
       'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'registered', 'cnt'],
      dtype='object')
```

```
In [18]: bike.nunique()
```

```
Out[18]: season        4
          yr           2
          mnth         12
          holiday      2
          weekday      7
          workingday   2
          weathersit    3
          temp          498
          atemp         689
          hum           594
          windspeed     649
          registered    678
          cnt           695
          dtype: int64
```

```
In [19]: cols = ['temp', 'atemp', 'hum', 'windspeed']
plt.figure(figsize=(18,4))
```

```
i = 1
for col in cols:
    plt.subplot(1,4,i)
    sns.boxplot(y=col, data=bike)
    i+=1
```



So on above plots we don't have any outliers

EDA

```
In [20]: #changing categorical data which were primarily numeric to more meaningful one
bike.season.replace({1:"spring", 2:"summer", 3:"fall", 4:"winter"},inplace = True)

bike.weathersit.replace({1:'good',2:'moderate',3:'bad',4:'severe'},inplace = True)

bike.mnth = bike.mnth.replace({1: 'jan',2: 'feb',3: 'mar',4: 'apr',5: 'may',6: 'jun',
    7: 'jul',8: 'aug',9: 'sept',10: 'oct',11: 'nov',12: 'dec'})

bike.weekday = bike.weekday.replace({0: 'sun',1: 'mon',2: 'tue',3: 'wed',4: 'thu',5: 'fri',6: 'sat'})
bike.head()
```

Out[20]:

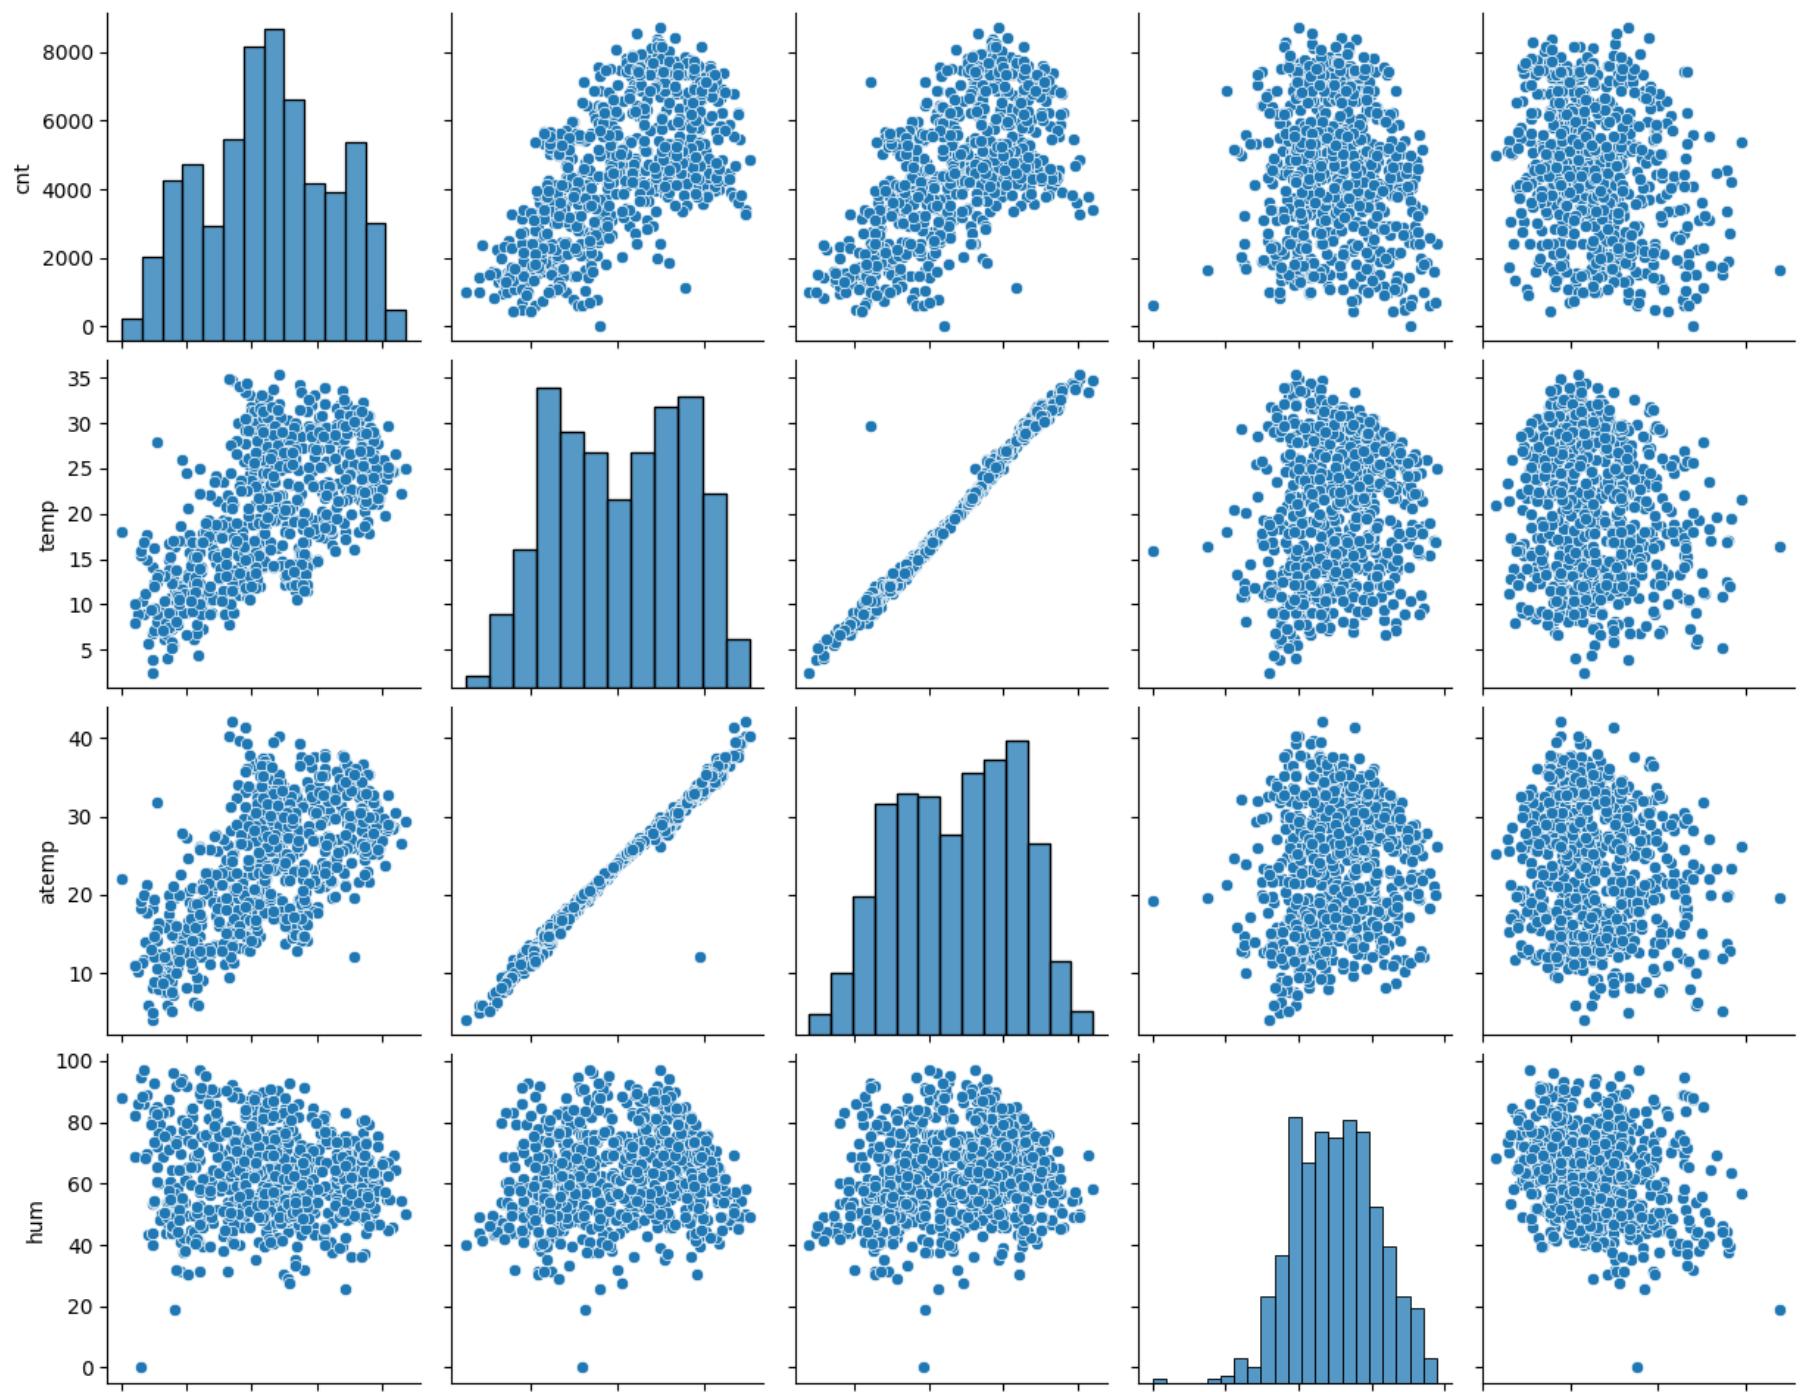
	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	registered	cnt
0	spring	0	jan	0	mon	1	moderate	14.110847	18.18125	80.5833	10.749882	654	985
1	spring	0	jan	0	tue	1	moderate	14.902598	17.68695	69.6087	16.652113	670	801
2	spring	0	jan	0	wed	1	good	8.050924	9.47025	43.7273	16.636703	1229	1349
3	spring	0	jan	0	thu	1	good	8.200000	10.60610	59.0435	10.739832	1454	1562
4	spring	0	jan	0	fri	1	good	9.305237	11.46350	43.6957	12.522300	1518	1600

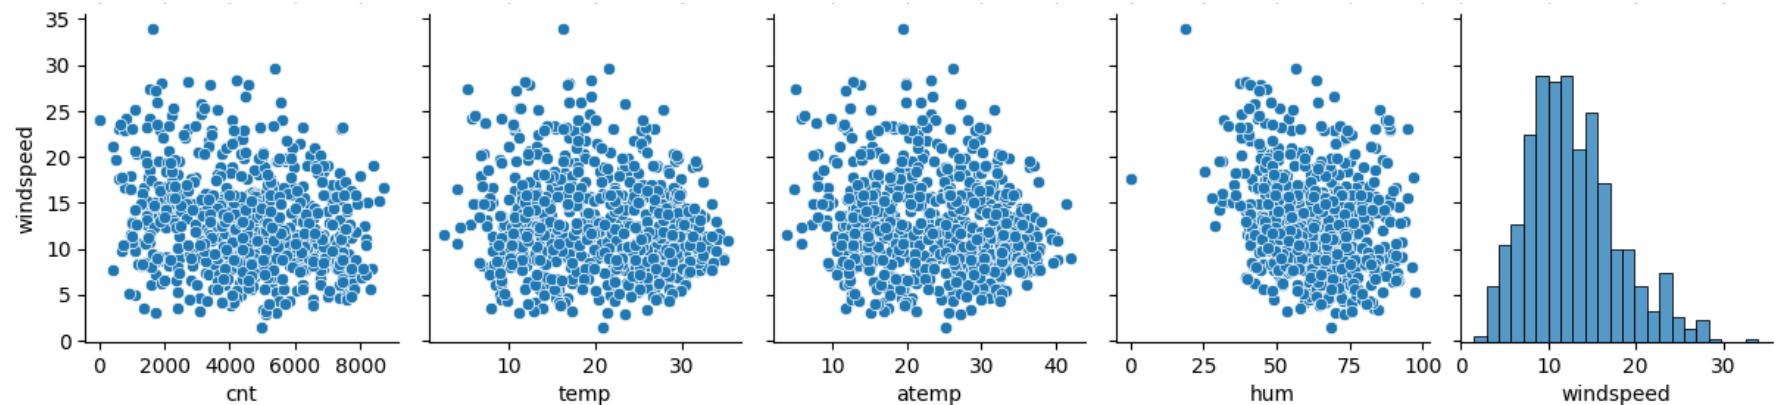
Drawing pairplots to check for linear relationship

In [21]:

```
plt.figure(figsize = (15,30))
sns.pairplot(data=bike,vars=['cnt', 'temp', 'atemp', 'hum','windspeed'])
plt.show()
```

<Figure size 1500x3000 with 0 Axes>





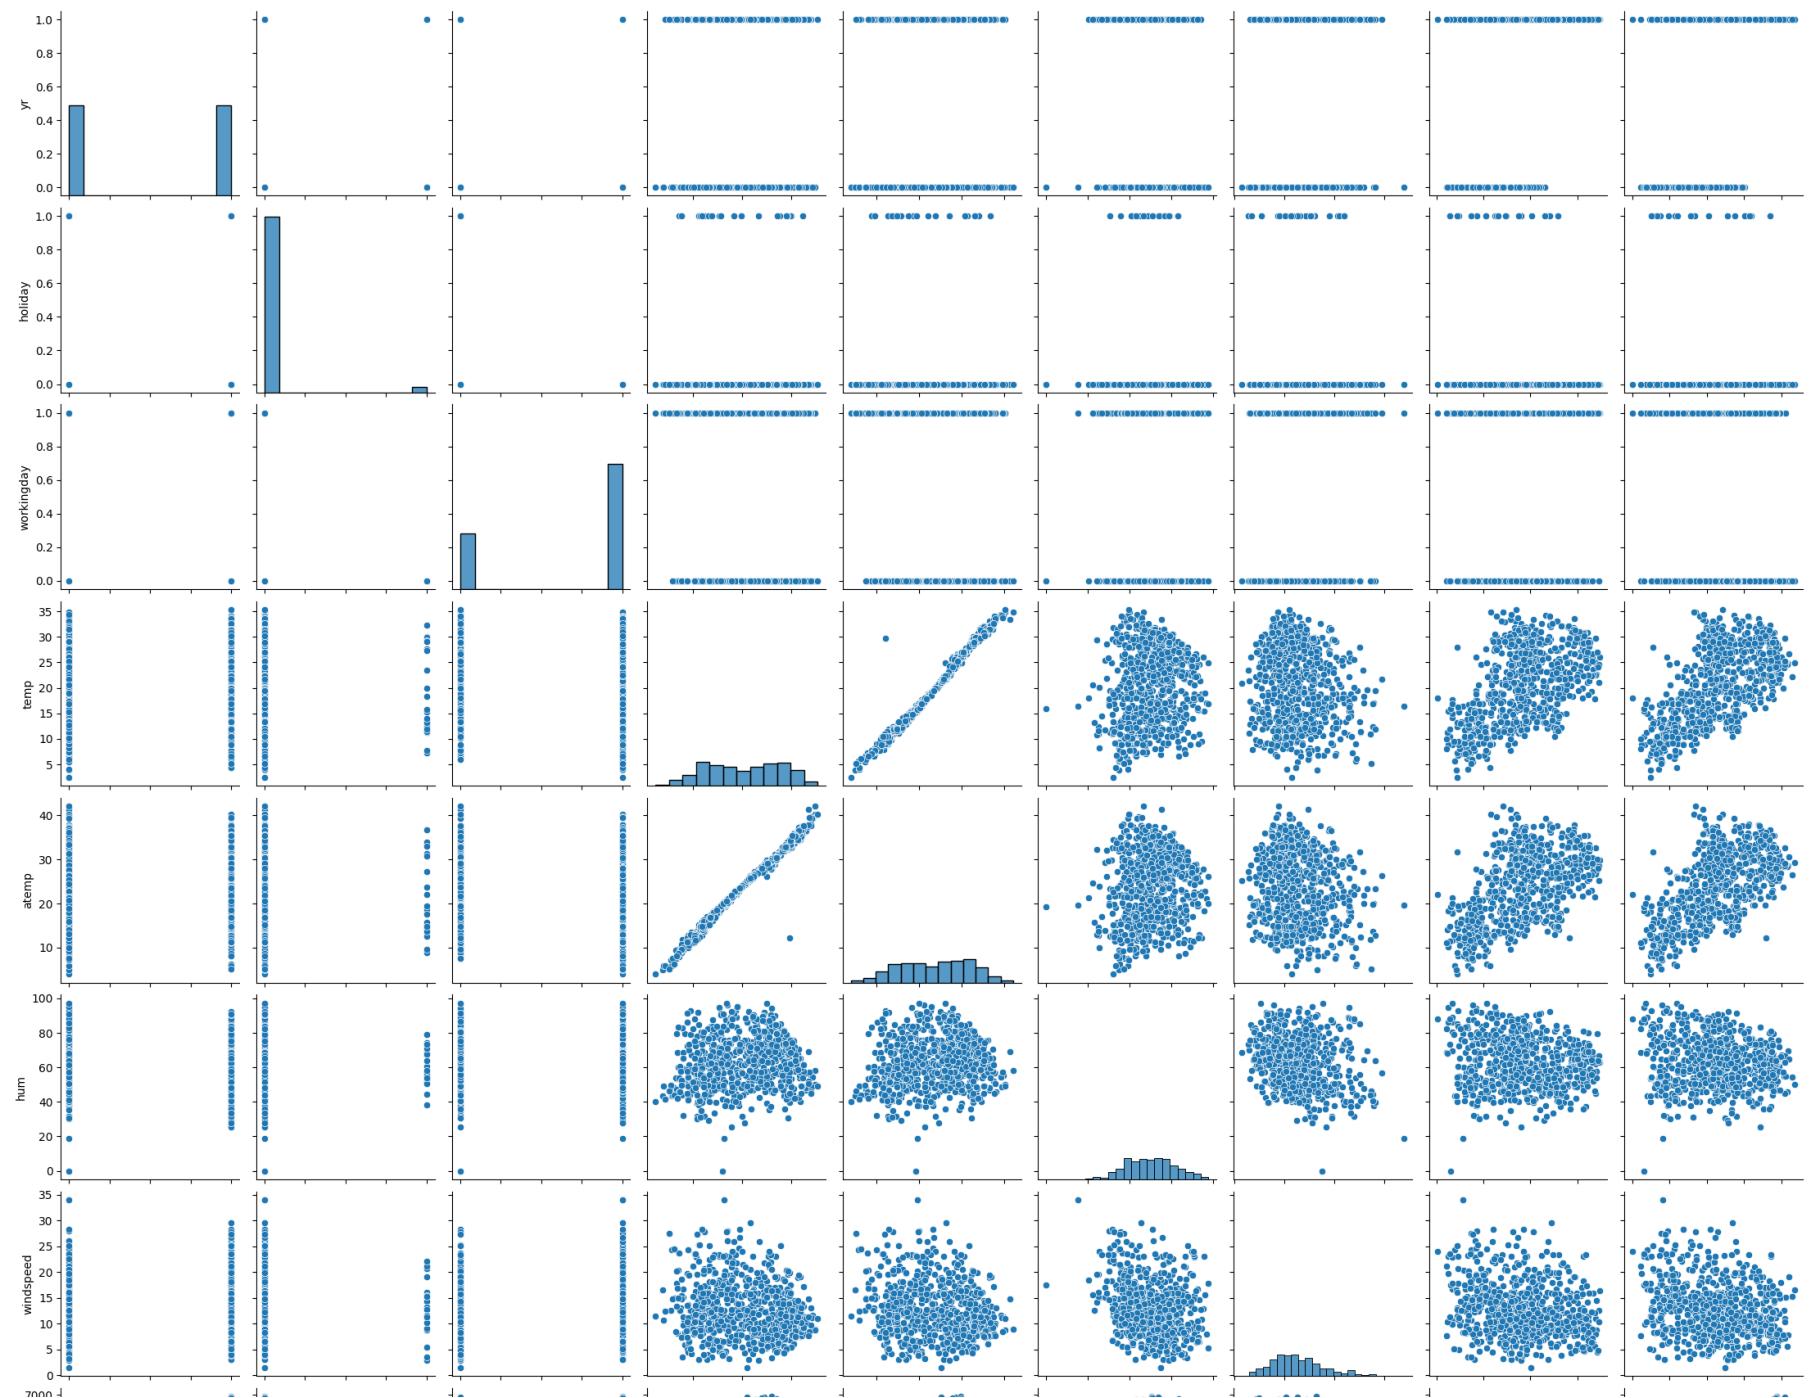
- So from the above plots we can clearly understand that temp and atemp are having high correlation
- And from the plots we can also say that there is a linear relationship between TEMP and ATEMP

Visualising the Data to find correlation from numerical variables

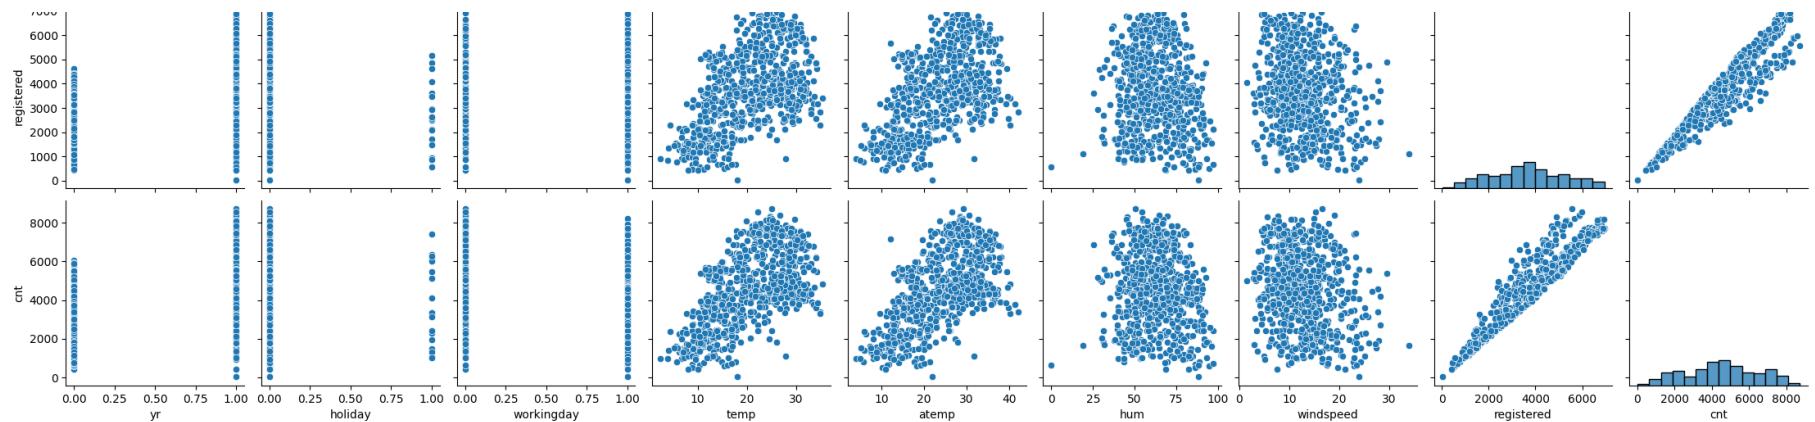
```
In [22]: plt.figure(figsize=(20,15))
sns.pairplot(bike)
plt.show()

<Figure size 2000x1500 with 0 Axes>
```

BikeShared



BikeShared



Heatmap for coorelation

```
In [23]: plt.figure(figsize=(15,10))
sns.heatmap(bike.corr(),cmap="BuPu", annot=True)
plt.show()
```



```
In [24]: #visualising the categorical variables

plt.figure(figsize=(20,15))

plt.subplot(2,3,1)
sns.boxplot(x='season',y='cnt',data=bike)

plt.subplot(2,3,2)
sns.boxplot(x='yr',y='cnt',data=bike)

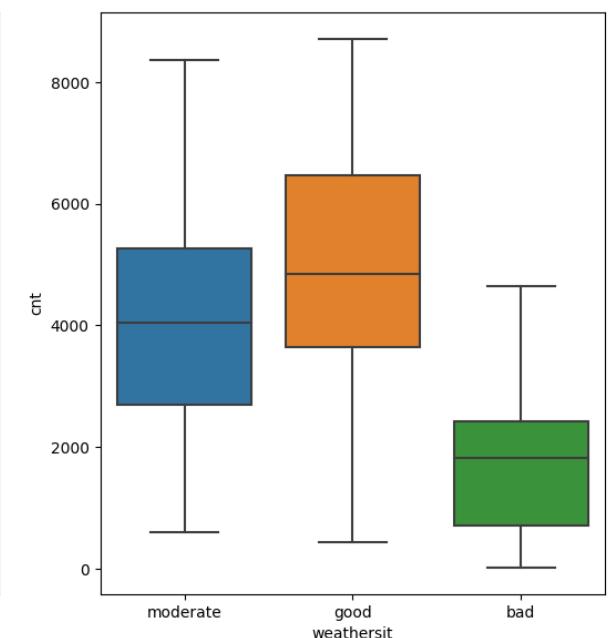
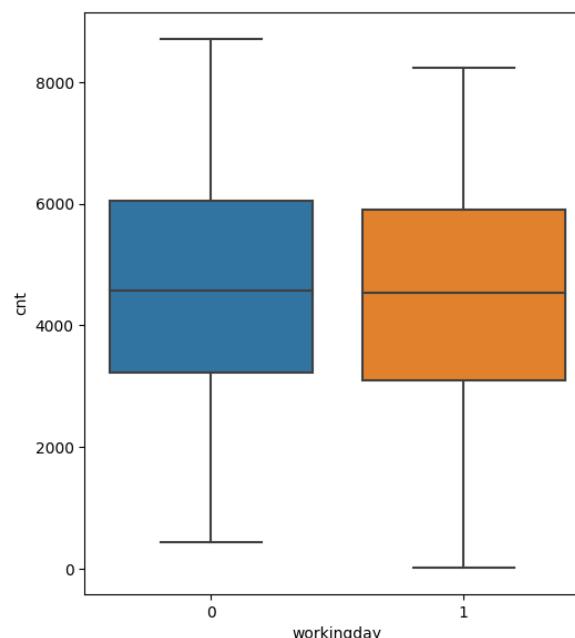
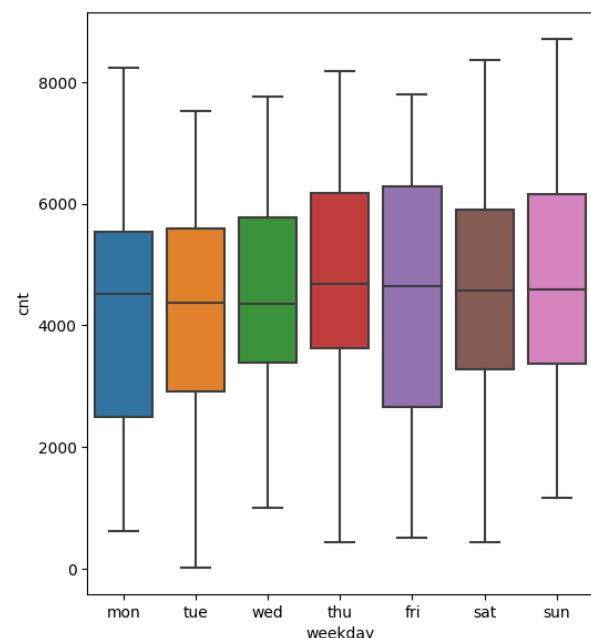
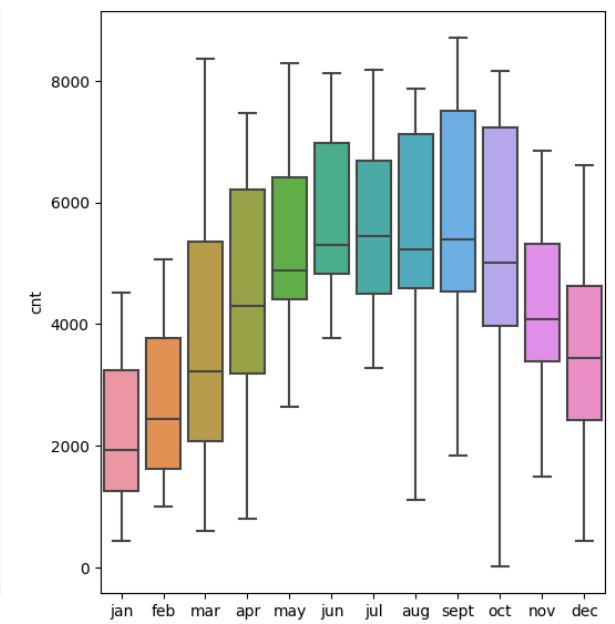
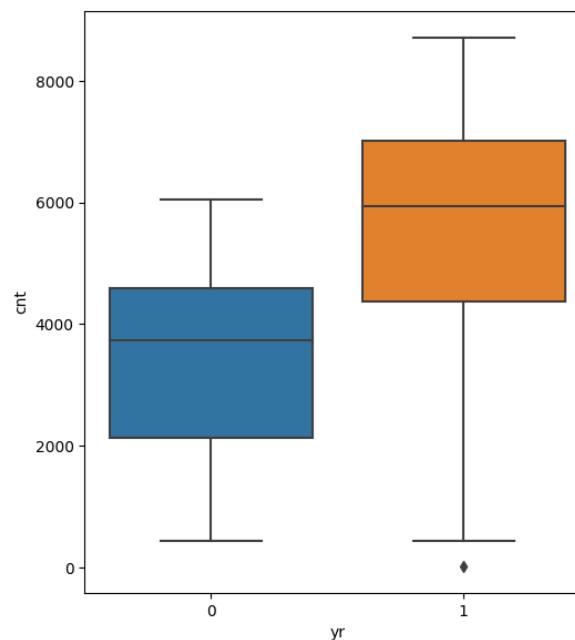
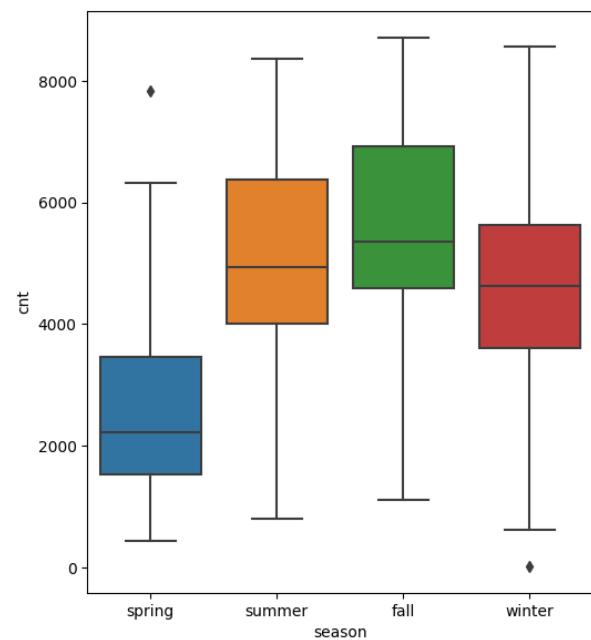
plt.subplot(2,3,3)
sns.boxplot(x='mnth',y='cnt',data=bike)

plt.subplot(2,3,4)
sns.boxplot(x='weekday',y='cnt',data=bike)

plt.subplot(2,3,5)
sns.boxplot(x='workingday',y='cnt',data=bike)

plt.subplot(2,3,6)
sns.boxplot(x='weathersit',y='cnt',data=bike)
```

```
Out[24]: <Axes: xlabel='weathersit', ylabel='cnt'>
```

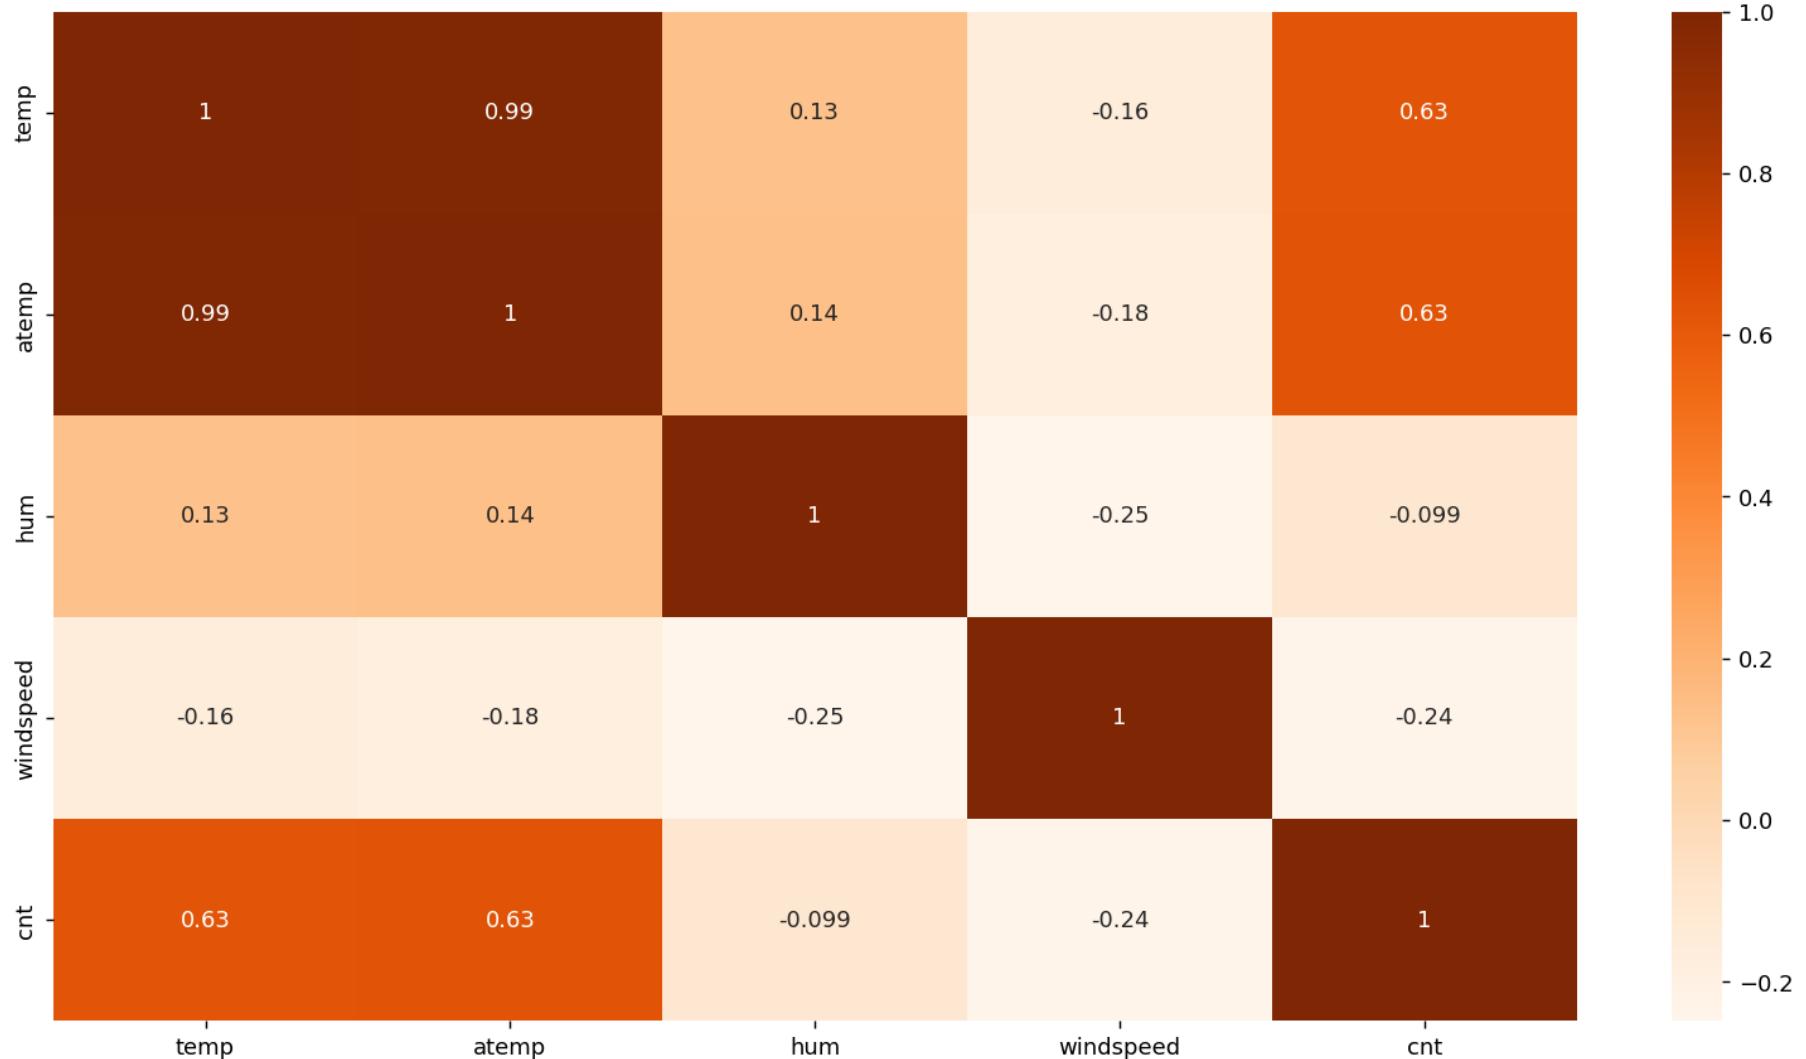


Conclusions:

- The graph clearly shows the qualitative distributions of the data, now if the model suggests the important predictors, using these graphs we can be more confident about the predictions of the model.
- For the variable season, we can clearly see that the category 3 : Fall, has the highest median, which shows that the demand was high during this season. It is least for 1: spring .
- The year 2019 had a higher count of users as compared to the year 2018
- The bike demand is almost constant throughout the week.
- The count of total users is in between 4000 to 6000 (~5500) during clear weather
- The count is highest in the month of August
- The count of users is less during the holidays

```
In [25]: num_features = ["temp", "atemp", "hum", "windspeed", "cnt"]
plt.figure(figsize=(15,8),dpi=130)
plt.title("Correlation between numeric features", fontsize=16)
sns.heatmap(bike[num_features].corr(), annot=True, cmap="Oranges")
plt.show()
```

Correlation betweenen numeric features

In [26]: `bike.describe()`

Out[26]:

	yr	holiday	workingday	temp	atemp	hum	windspeed	registered	cnt
count	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000
mean	0.500000	0.028767	0.690411	20.319259	23.726322	62.765175	12.763620	3658.757534	4508.006849
std	0.500343	0.167266	0.462641	7.506729	8.150308	14.237589	5.195841	1559.758728	1936.011647
min	0.000000	0.000000	0.000000	2.424346	3.953480	0.000000	1.500244	20.000000	22.000000
25%	0.000000	0.000000	0.000000	13.811885	16.889713	52.000000	9.041650	2502.250000	3169.750000
50%	0.500000	0.000000	1.000000	20.465826	24.368225	62.625000	12.125325	3664.500000	4548.500000
75%	1.000000	0.000000	1.000000	26.880615	30.445775	72.989575	15.625589	4783.250000	5966.000000
max	1.000000	1.000000	1.000000	35.328347	42.044800	97.250000	34.000021	6946.000000	8714.000000

Data preparation for linear regression:

Creating dummy variables for categorical variables:

```
In [27]: bike = pd.get_dummies(data=bike,columns=["season","mnth","weekday"],drop_first=True)
bike = pd.get_dummies(data=bike,columns=["weathersit"])
```

```
In [28]: bike.columns
```

```
Out[28]: Index(['yr', 'holiday', 'workingday', 'temp', 'atemp', 'hum', 'windspeed',
       'registered', 'cnt', 'season_spring', 'season_summer', 'season_winter',
       'mnth_aug', 'mnth_dec', 'mnth_feb', 'mnth_jan', 'mnth_jul', 'mnth_jun',
       'mnth_mar', 'mnth_may', 'mnth_nov', 'mnth_oct', 'mnth_sept',
       'weekday_mon', 'weekday_sat', 'weekday_sun', 'weekday_thu',
       'weekday_tue', 'weekday_wed', 'weathersit_bad', 'weathersit_good',
       'weathersit_moderate'],
      dtype='object')
```

```
In [29]: bike.head()
```

Out[29]:

	yr	holiday	workingday	temp	atemp	hum	windspeed	registered	cnt	season_spring	...	mnth_sept	weekday_mon	weekday_sat	...
0	0	0	1	14.110847	18.18125	80.5833	10.749882	654	985	1	...	0	1	1	0
1	0	0	1	14.902598	17.68695	69.6087	16.652113	670	801	1	...	0	0	0	0
2	0	0	1	8.050924	9.47025	43.7273	16.636703	1229	1349	1	...	0	0	0	0
3	0	0	1	8.200000	10.60610	59.0435	10.739832	1454	1562	1	...	0	0	0	0
4	0	0	1	9.305237	11.46350	43.6957	12.522300	1518	1600	1	...	0	0	0	0

5 rows × 32 columns

Splitting Data into train and test data:

In [30]: `bike.shape`

Out[30]: (730, 32)

In [31]:

```
#y to contain only target variable
y=bike.pop('cnt')

#X is all remainign variable also our independent variables
X=bike

#Train Test split with 70:30 ratio
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

In [32]: `X.head()`

Out[32]:

	yr	holiday	workingday	temp	atemp	hum	windspeed	registered	season_spring	season_summer	...	mnth_sept	weekday_mon	weel
0	0	0	1	14.110847	18.18125	80.5833	10.749882	654	1	0	...	0	0	1
1	0	0	1	14.902598	17.68695	69.6087	16.652113	670	1	0	...	0	0	0
2	0	0	1	8.050924	9.47025	43.7273	16.636703	1229	1	0	...	0	0	0
3	0	0	1	8.200000	10.60610	59.0435	10.739832	1454	1	0	...	0	0	0
4	0	0	1	9.305237	11.46350	43.6957	12.522300	1518	1	0	...	0	0	0

5 rows × 31 columns

In [33]:

```
print(X_train.shape)
print(X_test.shape)
```

(511, 31)
(219, 31)

In [34]:

```
num_vars = ['temp', 'atemp', 'hum', 'windspeed']
```

```
#Use Normalized scaler to scale
scaler = MinMaxScaler()

#Fit and transform training set only
X_train[num_vars] = scaler.fit_transform(X_train[num_vars])
```

In [35]:

```
X_train.describe()
```

Out[35]:

	yr	holiday	workingday	temp	atemp	hum	windspeed	registered	season_spring	season_summer	...	mnth_se
count	511.000000	511.000000	511.000000	511.000000	511.000000	511.000000	511.000000	511.000000	511.000000	511.000000	...	511.000000
mean	0.520548	0.027397	0.698630	0.550874	0.526518	0.654174	0.414873	3694.350294	0.230920	0.258317	...	0.090
std	0.500067	0.163398	0.459303	0.227231	0.212976	0.143648	0.185619	1573.398840	0.421834	0.438138	...	0.286
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	416.000000	0.000000	0.000000	...	0.000
25%	0.000000	0.000000	0.000000	0.354221	0.353207	0.544598	0.281934	2628.000000	0.000000	0.000000	...	0.000
50%	1.000000	0.000000	1.000000	0.575304	0.563297	0.659940	0.384732	3729.000000	0.000000	0.000000	...	0.000
75%	1.000000	0.000000	1.000000	0.745079	0.698389	0.756977	0.526481	4804.000000	0.000000	1.000000	...	0.000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	6946.000000	1.000000	1.000000	...	1.000

8 rows × 31 columns

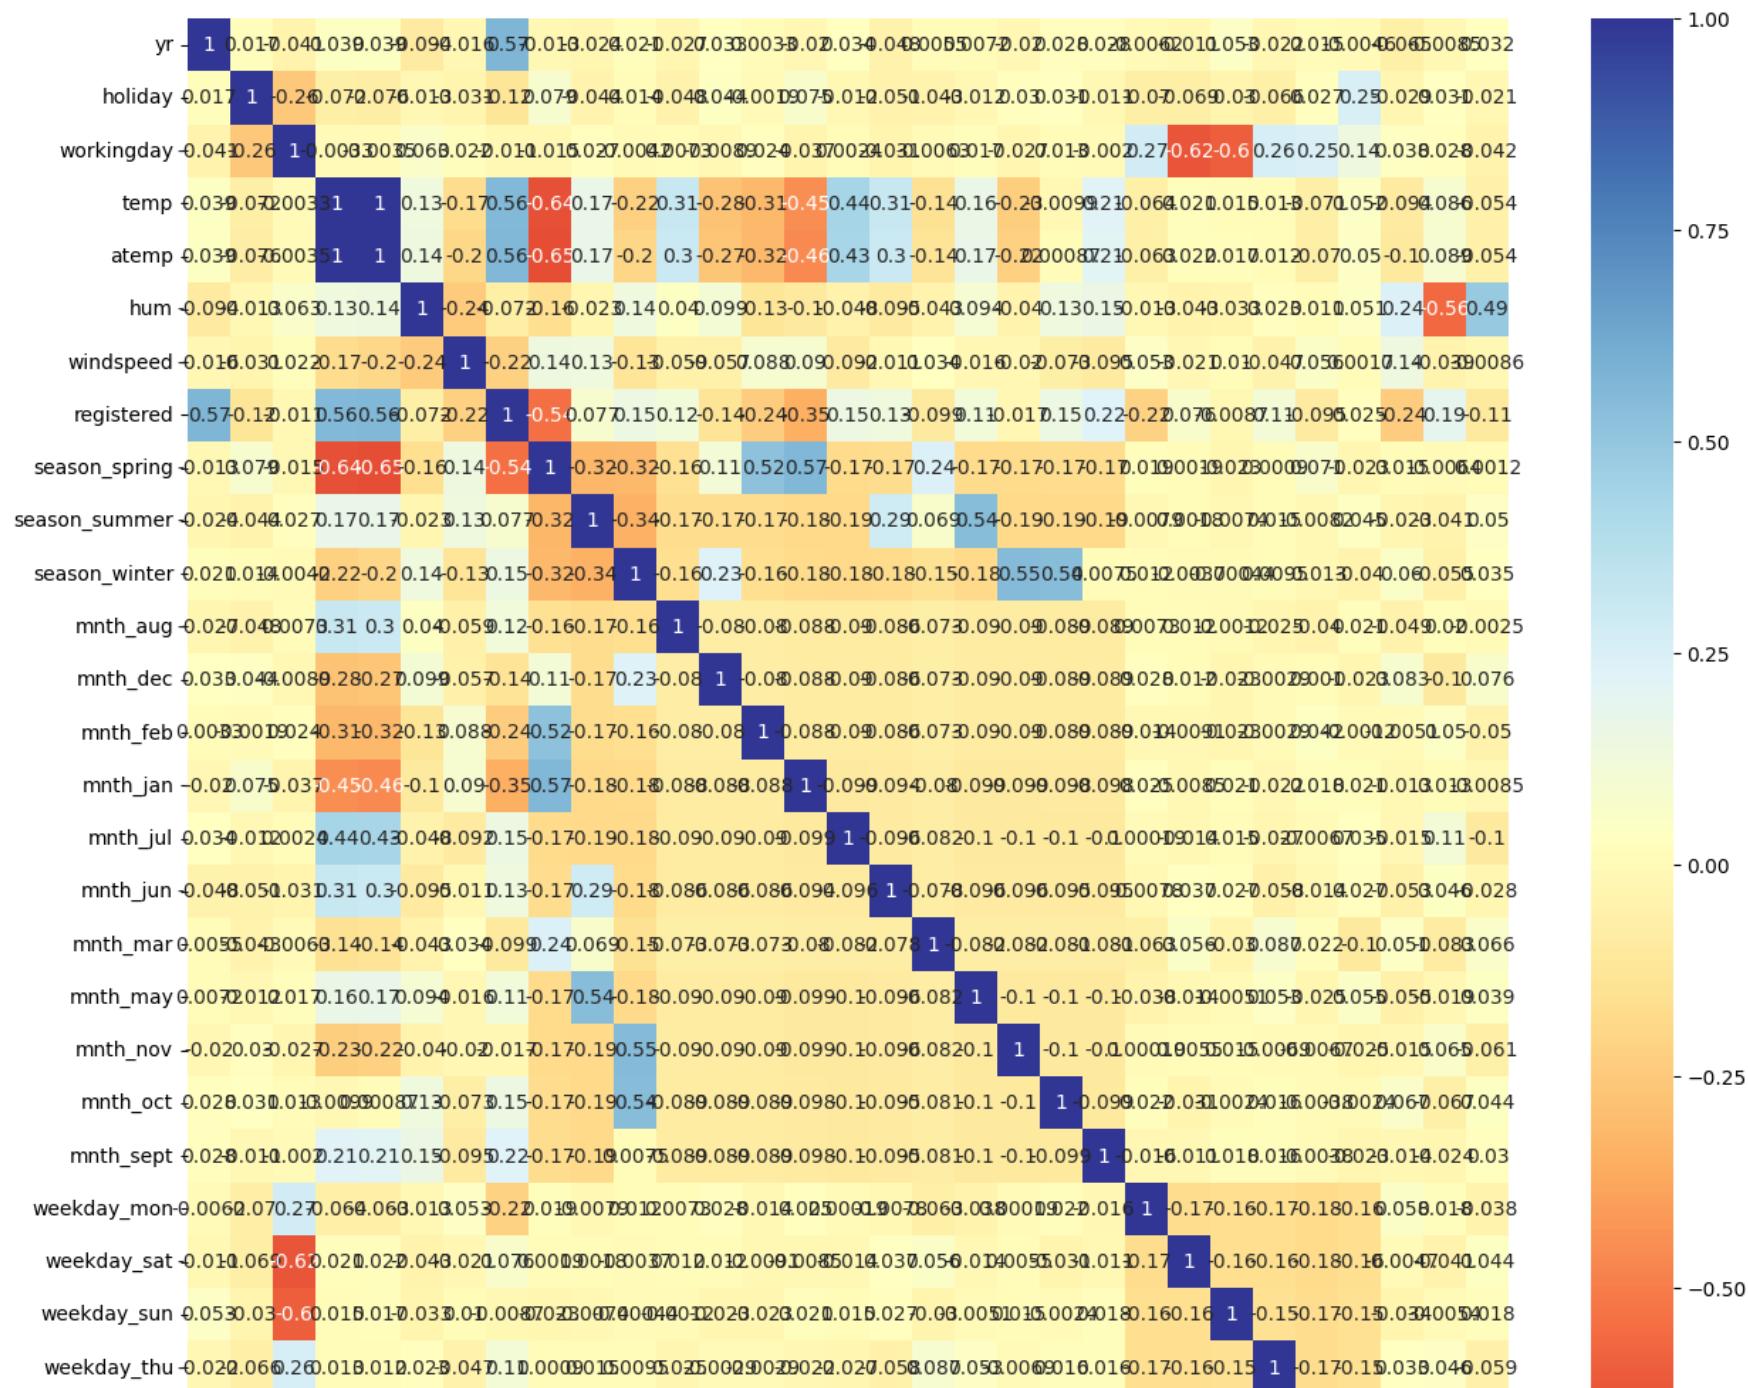
In [36]: `x_train.head()`

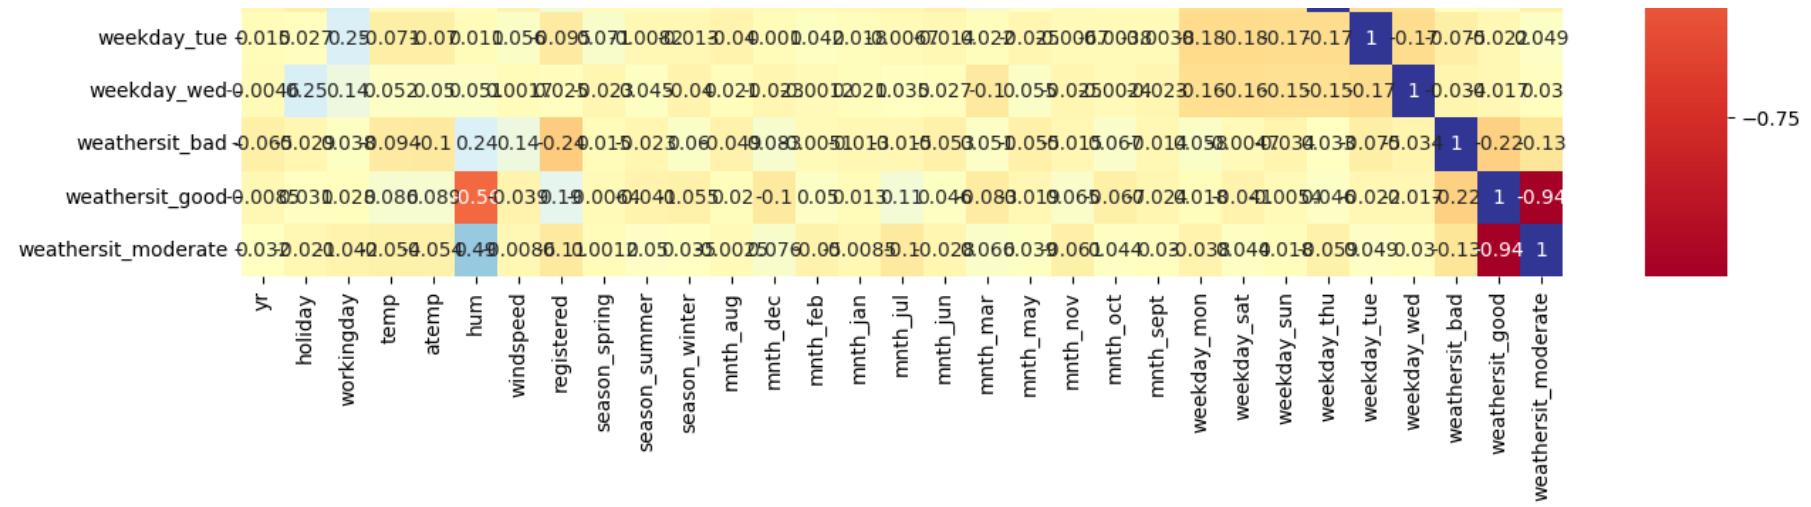
Out[36]:

	yr	holiday	workingday	temp	atemp	hum	windspeed	registered	season_spring	season_summer	...	mnth_sept	weekday_mon	we
683	1	0	1	0.327133	0.322150	0.639330	0.327101	5125	0	0	...	0	0	0
645	1	1	0	0.403972	0.404998	0.731215	0.419004	4604	0	0	...	0	0	0
163	0	0	1	0.717562	0.685963	0.509660	0.708724	4157	0	1	...	0	0	0
360	0	0	1	0.331287	0.326273	0.785745	0.415925	1059	1	0	...	0	0	0
640	1	0	1	0.745598	0.682653	0.817947	0.110593	6844	0	0	...	0	0	0

5 rows × 31 columns

In [37]: `plt.figure(figsize = (15, 15)) #Checking if the variables are highly correlated
sns.heatmap(X_train.corr(), annot = True, cmap='RdYlBu')
plt.show()`





Build model using RFE

```
In [38]: lr = LinearRegression()
lr.fit(X_train,y_train)
```

Out[38]:

LinearRegression

LinearRegression()

```
In [39]: # Create the estimator (Linear Regression in this case)
lr = LinearRegression()

# Create the RFE object with the estimator and the number of features to select (15 in this case)
rfe = RFE(estimator=lr, n_features_to_select=15)

# Fit the RFE object to the training data
rfe.fit(X_train, y_train)
```

Out[39]:

```
▶ RFE
  ▶ estimator: LinearRegression
    ▶ LinearRegression
```

In [40]: `list(zip(X_train.columns,rfe.support_,rfe.ranking_))`

```
Out[40]: [('yr', True, 1),
           ('holiday', True, 1),
           ('workingday', False, 13),
           ('temp', True, 1),
           ('atemp', True, 1),
           ('hum', True, 1),
           ('windspeed', True, 1),
           ('registered', False, 17),
           ('season_spring', True, 1),
           ('season_summer', False, 8),
           ('season_winter', True, 1),
           ('mnth_aug', False, 12),
           ('mnth_dec', True, 1),
           ('mnth_feb', False, 9),
           ('mnth_jan', False, 15),
           ('mnth_jul', True, 1),
           ('mnth_jun', False, 11),
           ('mnth_mar', True, 1),
           ('mnth_may', False, 2),
           ('mnth_nov', True, 1),
           ('mnth_oct', False, 4),
           ('mnth_sept', True, 1),
           ('weekday_mon', False, 5),
           ('weekday_sat', False, 14),
           ('weekday_sun', False, 10),
           ('weekday_thu', False, 16),
           ('weekday_tue', False, 7),
           ('weekday_wed', False, 6),
           ('weathersit_bad', True, 1),
           ('weathersit_good', True, 1),
           ('weathersit_moderate', False, 3)]
```

```
In [41]: def build_model(cols):
    X_train_sm = sm.add_constant(X_train[cols])
    lm = sm.OLS(y_train, X_train_sm).fit()
    print(lm.summary())
    return lm
```

```
In [42]: def get_vif(cols):
    df1 = X_train[cols]
    vif = pd.DataFrame()
    vif['Features'] = df1.columns
    vif['VIF'] = [variance_inflation_factor(df1.values, i) for i in range(df1.shape[1])]
    vif['VIF'] = round(vif['VIF'],2)
    print(vif.sort_values(by='VIF', ascending=False))
```

```
In [43]: #Print Columns selected by RFE. We will manually eliminate for these columns
X_train.columns[rfe.support_]
```

```
Out[43]: Index(['yr', 'holiday', 'temp', 'atemp', 'hum', 'windspeed', 'season_spring',
       'season_winter', 'mnth_dec', 'mnth_jul', 'mnth_mar', 'mnth_nov',
       'mnth_sept', 'weathersit_bad', 'weathersit_good'],
      dtype='object')
```

```
In [44]: # Features not selected by RFE
X_train.columns[~rfe.support_]
```

```
Out[44]: Index(['workingday', 'registered', 'season_summer', 'mnth_aug', 'mnth_feb',
       'mnth_jan', 'mnth_jun', 'mnth_may', 'mnth_oct', 'weekday_mon',
       'weekday_sat', 'weekday_sun', 'weekday_thu', 'weekday_tue',
       'weekday_wed', 'weathersit_moderate'],
      dtype='object')
```

```
In [45]: # Taking 15 columns supported by RFE for regression
X_train_rfe = X_train[['yr', 'holiday', 'workingday', 'temp', 'hum', 'windspeed', 'season_spring',
       'season_summer', 'season_winter', 'mnth_jan', 'mnth_jul', 'mnth_sept', 'weekday_sat',
       'weathersit_bad', 'weathersit_moderate']]
```

```
In [46]: X_train_rfe.shape
```

Out[46]: (511, 15)

Model 1

```
In [47]: #Selected columns for Model 1 - all columns selected by RFE
cols = ['yr', 'holiday', 'workingday', 'temp', 'hum', 'windspeed', 'season_spring',
        'season_summer', 'season_winter', 'mnth_jan', 'mnth_jul', 'mnth_sept', 'weekday_sat',
        'weathersit_bad', 'weathersit_moderate']

build_model(cols)
get_vif(cols)
```

OLS Regression Results

Dep. Variable:	cnt	R-squared:	0.833			
Model:	OLS	Adj. R-squared:	0.828			
Method:	Least Squares	F-statistic:	165.0			
Date:	Tue, 01 Aug 2023	Prob (F-statistic):	1.70e-181			
Time:	13:30:57	Log-Likelihood:	-4134.9			
No. Observations:	511	AIC:	8302.			
Df Residuals:	495	BIC:	8370.			
Df Model:	15					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	2809.2461	346.564	8.106	0.000	2128.328	3490.164
yr	1955.0549	72.324	27.032	0.000	1812.955	2097.155
holiday	-839.6361	237.519	-3.535	0.000	-1306.306	-372.966
workingday	-202.7038	107.596	-1.884	0.060	-414.104	8.697
temp	4233.9838	303.827	13.936	0.000	3637.034	4830.934
hum	-1378.2895	339.821	-4.056	0.000	-2045.959	-710.620
windspeed	-1072.6371	210.709	-5.091	0.000	-1486.631	-658.643
season_spring	-795.1544	191.300	-4.157	0.000	-1171.014	-419.294
season_summer	380.6325	133.154	2.859	0.004	119.015	642.250
season_winter	767.1240	151.761	5.055	0.000	468.950	1065.298
mnth_jan	-59.8050	156.115	-0.383	0.702	-366.534	246.924
mnth_jul	-367.2369	159.579	-2.301	0.022	-680.772	-53.702
mnth_sept	666.1046	144.332	4.615	0.000	382.526	949.684
weekday_sat	-72.5004	136.440	-0.531	0.595	-340.574	195.573
weathersit_bad	-1761.9834	238.533	-7.387	0.000	-2230.645	-1293.322
weathersit_moderate	-464.2948	93.779	-4.951	0.000	-648.550	-280.040
Omnibus:	65.698	Durbin-Watson:	2.095			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	125.994			
Skew:	-0.750	Prob(JB):	4.37e-28			
Kurtosis:	4.914	Cond. No.	23.6			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Features VIF

```
4          hum  28.53
3        temp  17.32
2   workingday  6.06
5      windspeed  5.73
6  season_spring  4.45
8  season_winter  3.24
7  season_summer  3.01
14 weathersit_moderate  2.26
0           yr  2.10
12 weekday_sat  2.01
10 mnth_jul  1.83
9  mnth_jan  1.67
11  mnth_sept  1.47
13 weathersit_bad  1.27
1       holiday  1.21
```

Model 2

```
In [48]: # Dropping the variable mnth_jan as it has negative coefficient and is insignificant as it has high p-value
cols = ['yr', 'holiday', 'workingday', 'temp', 'hum', 'windspeed', 'season_spring',
        'season_summer', 'season_winter', 'mnth_jul', 'mnth_sept', 'weekday_sat',
        'weathersit_bad', 'weathersit_moderate']
build_model(cols)
get_vif(cols)
```

OLS Regression Results

Dep. Variable:	cnt	R-squared:	0.833			
Model:	OLS	Adj. R-squared:	0.829			
Method:	Least Squares	F-statistic:	177.1			
Date:	Tue, 01 Aug 2023	Prob (F-statistic):	1.35e-182			
Time:	13:31:02	Log-Likelihood:	-4135.0			
No. Observations:	511	AIC:	8300.			
Df Residuals:	496	BIC:	8363.			
Df Model:	14					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	2792.2799	343.427	8.131	0.000	2117.529	3467.030
yr	1955.1531	72.261	27.057	0.000	1813.177	2097.129
holiday	-840.8434	237.294	-3.543	0.000	-1307.069	-374.618
workingday	-200.8684	107.396	-1.870	0.062	-411.877	10.140
temp	4258.5248	296.741	14.351	0.000	3675.500	4841.549
hum	-1383.9501	339.207	-4.080	0.000	-2050.411	-717.489
windspeed	-1073.0058	210.525	-5.097	0.000	-1486.637	-659.375
season_spring	-806.5438	188.813	-4.272	0.000	-1177.516	-435.572
season_summer	384.0827	132.735	2.894	0.004	123.290	644.875
season_winter	774.4415	150.424	5.148	0.000	478.895	1069.988
mnth_jul	-369.6007	159.322	-2.320	0.021	-682.630	-56.571
mnth_sept	665.9059	144.207	4.618	0.000	382.574	949.238
weekday_sat	-71.1261	136.276	-0.522	0.602	-338.875	196.623
weathersit_bad	-1756.5848	237.911	-7.383	0.000	-2224.023	-1289.147
weathersit_moderate	-462.7865	93.616	-4.943	0.000	-646.719	-278.854
Omnibus:	65.111	Durbin-Watson:	2.095			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	124.284			
Skew:	-0.746	Prob(JB):	1.03e-27			
Kurtosis:	4.900	Cond. No.	23.5			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Features	VIF
4	hum 27.99

```
3          temp  16.83
2      workingday   6.06
5      windspeed   5.70
6  season_spring  4.10
8  season_winter  3.23
7  season_summer  3.01
13 weathersit_moderate  2.25
0          yr    2.10
11 weekday_sat   2.01
9      mnth_jul   1.82
10     mnth_sept  1.47
12 weathersit_bad  1.26
1      holiday    1.21
```

Model-3

```
In [49]: # Dropping the variable hum as it has negative coefficient and is insignificant as it has high p-value
cols = ['yr', 'holiday', 'workingday', 'temp', 'windspeed', 'season_spring',
        'season_summer', 'season_winter', 'mnth_jul', 'mnth_sept', 'weekday_sat',
        'weathersit_bad', 'weathersit_moderate']
build_model(cols)
get_vif(cols)
```

OLS Regression Results

=====						
Dep. Variable:	cnt	R-squared:	0.828			
Model:	OLS	Adj. R-squared:	0.823			
Method:	Least Squares	F-statistic:	183.6			
Date:	Tue, 01 Aug 2023	Prob (F-statistic):	3.37e-180			
Time:	13:31:06	Log-Likelihood:	-4143.4			
No. Observations:	511	AIC:	8315.			
Df Residuals:	497	BIC:	8374.			
Df Model:	13					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	2004.0762	288.371	6.950	0.000	1437.500	2570.652
yr	1995.7722	72.690	27.456	0.000	1852.954	2138.590
holiday	-873.2715	240.865	-3.626	0.000	-1346.511	-400.032
workingday	-226.4412	108.888	-2.080	0.038	-440.379	-12.504
temp	4050.9306	296.912	13.644	0.000	3467.573	4634.289
windspeed	-816.0345	204.020	-4.000	0.000	-1216.883	-415.186
season_spring	-838.0753	191.601	-4.374	0.000	-1214.523	-461.627
season_summer	376.3936	134.795	2.792	0.005	111.556	641.231
season_winter	713.5094	152.019	4.694	0.000	414.831	1012.188
mnth_jul	-326.5137	161.455	-2.022	0.044	-643.731	-9.296
mnth_sept	610.3106	145.804	4.186	0.000	323.843	896.779
weekday_sat	-50.4406	138.309	-0.365	0.715	-322.182	221.301
weathersit_bad	-2161.8583	219.559	-9.846	0.000	-2593.236	-1730.481
weathersit_moderate	-685.8621	77.177	-8.887	0.000	-837.495	-534.229
=====						
Omnibus:	60.891	Durbin-Watson:	2.127			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	113.587			
Skew:	-0.712	Prob(JB):	2.16e-25			
Kurtosis:	4.819	Cond. No.	20.3			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

	Features	VIF
3	temp	8.10
2	workingday	5.73

```
4      windspeed  5.69
5      season_spring  2.69
6      season_summer  2.69
7      season_winter  2.19
0          yr  2.09
10     weekday_sat  1.99
8          mnth_jul  1.82
12  weathersit_moderate  1.57
9          mnth_sept  1.43
1      holiday  1.19
11  weathersit_bad  1.09
```

Model-4

```
In [50]: # Dropping the variable holiday as it has negative coefficient and is insignificant as it has high p-value
cols = ['yr', 'workingday', 'temp', 'windspeed', 'season_spring',
        'season_summer', 'season_winter', 'mnth_jul', 'mnth_sept', 'weekday_sat',
        'weathersit_bad', 'weathersit_moderate']
build_model(cols)
get_vif(cols)
```

OLS Regression Results

=====						
Dep. Variable:	cnt	R-squared:	0.823			
Model:	OLS	Adj. R-squared:	0.819			
Method:	Least Squares	F-statistic:	193.1			
Date:	Tue, 01 Aug 2023	Prob (F-statistic):	1.51e-178			
Time:	13:31:09	Log-Likelihood:	-4150.1			
No. Observations:	511	AIC:	8326.			
Df Residuals:	498	BIC:	8381.			
Df Model:	12					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	1845.8415	288.504	6.398	0.000	1279.007	2412.676
yr	1997.7265	73.569	27.154	0.000	1853.182	2142.271
workingday	-76.1270	101.906	-0.747	0.455	-276.346	124.093
temp	4062.1670	300.495	13.518	0.000	3471.774	4652.560
windspeed	-786.5884	206.329	-3.812	0.000	-1191.972	-381.205
season_spring	-870.5177	193.712	-4.494	0.000	-1251.111	-489.925
season_summer	366.4796	136.400	2.687	0.007	98.489	634.471
season_winter	697.9586	153.800	4.538	0.000	395.782	1000.135
mnth_jul	-335.0219	163.394	-2.050	0.041	-656.049	-13.995
mnth_sept	606.0378	147.566	4.107	0.000	316.109	895.967
weekday_sat	99.5698	133.574	0.745	0.456	-162.869	362.008
weathersit_bad	-2148.1594	222.187	-9.668	0.000	-2584.699	-1711.620
weathersit_moderate	-676.7543	78.071	-8.668	0.000	-830.143	-523.365
=====						
Omnibus:	68.186	Durbin-Watson:	2.118			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	126.740			
Skew:	-0.787	Prob(JB):	3.01e-28			
Kurtosis:	4.864	Cond. No.	20.2			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

	Features	VIF
2	temp	7.79
3	windspeed	5.69
1	workingday	5.01

```
5      season_summer 2.66
4      season_spring 2.51
6      season_winter 2.10
0          yr 2.09
9      weekday_sat 1.84
7          mnth_jul 1.82
11 weathersit_moderate 1.57
8          mnth_sept 1.42
10 weathersit_bad 1.09
```

Model-5

```
In [51]: # Dropping the variable mnth_jul,temp as it has negative coefficient and is insignificant as it has high p-value
cols = ['yr', 'workingday', 'temp', 'windspeed', 'season_spring',
        'season_summer', 'season_winter', 'mnth_sept', 'weekday_sat',
        'weathersit_bad', 'weathersit_moderate']
build_model(cols)
get_vif(cols)
```

OLS Regression Results

=====						
Dep. Variable:	cnt	R-squared:	0.822			
Model:	OLS	Adj. R-squared:	0.818			
Method:	Least Squares	F-statistic:	209.0			
Date:	Tue, 01 Aug 2023	Prob (F-statistic):	8.20e-179			
Time:	13:31:13	Log-Likelihood:	-4152.2			
No. Observations:	511	AIC:	8328.			
Df Residuals:	499	BIC:	8379.			
Df Model:	11					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
const	1779.2943	287.591	6.187	0.000	1214.255	2344.333
yr	1992.6143	73.763	27.014	0.000	1847.691	2137.538
workingday	-74.8818	102.231	-0.732	0.464	-275.739	125.975
temp	3958.5476	297.164	13.321	0.000	3374.701	4542.395
windspeed	-779.1681	206.959	-3.765	0.000	-1185.786	-372.550
season_spring	-778.9193	189.094	-4.119	0.000	-1150.438	-407.400
season_summer	491.7101	122.353	4.019	0.000	251.320	732.100
season_winter	798.5462	146.233	5.461	0.000	511.238	1085.855
mnth_sept	715.4337	138.024	5.183	0.000	444.253	986.614
weekday_sat	106.3578	133.961	0.794	0.428	-156.840	369.555
weathersit_bad	-2160.9326	222.811	-9.698	0.000	-2598.697	-1723.169
weathersit_moderate	-669.8623	78.249	-8.561	0.000	-823.600	-516.125
=====						
Omnibus:	70.746	Durbin-Watson:	2.113			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	130.655			
Skew:	-0.816	Prob(JB):	4.25e-29			
Kurtosis:	4.864	Cond. No.	20.2			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

	Features	VIF
2	temp	6.13
3	windspeed	5.68
1	workingday	5.01
4	season_spring	2.38

```
5      season_summer 2.10
0          yr 2.09
6      season_winter 1.85
8      weekday_sat 1.84
10 weathersit_moderate 1.57
7      mnth_sept 1.25
9      weathersit_bad 1.09
```

Model-6

```
In [52]: # Droping and replacing Variable as follows weekday_saturdat as high P-value, reconsidering holiday.
cols6 = ['yr','season_spring',
         'season_winter','season_summer','mnth_sept','holiday','workingday',
         'weathersit_bad', 'weathersit_moderate','temp']
build_model(cols6)
get_vif(cols6)
```

OLS Regression Results

=====						
Dep. Variable:	cnt	R-squared:	0.821			
Model:	OLS	Adj. R-squared:	0.817			
Method:	Least Squares	F-statistic:	229.0			
Date:	Tue, 01 Aug 2023	Prob (F-statistic):	1.70e-179			
Time:	13:31:18	Log-Likelihood:	-4153.4			
No. Observations:	511	AIC:	8329.			
Df Residuals:	500	BIC:	8375.			
Df Model:	10					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
const	1500.5024	262.410	5.718	0.000	984.940	2016.065
yr	1989.4538	73.778	26.965	0.000	1844.500	2134.408
season_spring	-744.1779	189.384	-3.929	0.000	-1116.264	-372.092
season_winter	865.5108	145.733	5.939	0.000	579.187	1151.834
season_summer	462.1289	122.140	3.784	0.000	222.159	702.099
mnth_sept	733.2217	138.143	5.308	0.000	461.809	1004.634
holiday	-825.2345	233.487	-3.534	0.000	-1283.970	-366.499
workingday	-204.1270	82.866	-2.463	0.014	-366.935	-41.319
weathersit_bad	-2289.0675	221.187	-10.349	0.000	-2723.638	-1854.497
weathersit_moderate	-678.4237	78.385	-8.655	0.000	-832.429	-524.419
temp	4091.5171	295.064	13.867	0.000	3511.798	4671.236
=====						
Omnibus:	63.271	Durbin-Watson:	2.151			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	118.647			
Skew:	-0.734	Prob(JB):	1.72e-26			
Kurtosis:	4.848	Cond. No.	19.0			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

	Features	VIF
9	temp	3.98
6	workingday	3.35
0	yr	2.08
3	season_summer	1.80
1	season_spring	1.62

```
2      season_winter  1.60
8  weathersit_moderate  1.56
4      mnth_sept   1.25
5      holiday    1.10
7  weathersit_bad   1.06
```

Here the VIF seems to be almost accepted. p-value for all the features is almost 0.0 and R2 is 0.8
Hence we finalize this model to use further

```
In [53]: #Build a model with all columns to select features automatically
def build_model_sk(X,y):
    lr1 = LinearRegression()
    lr1.fit(X,y)
    return lr1
```

```
In [54]: #Let us build the finalmodel using sklearn
#Build a model with above columns
lr = build_model_sk(X_train[cols6],y_train)
print(lr.intercept_,lr.coef_)

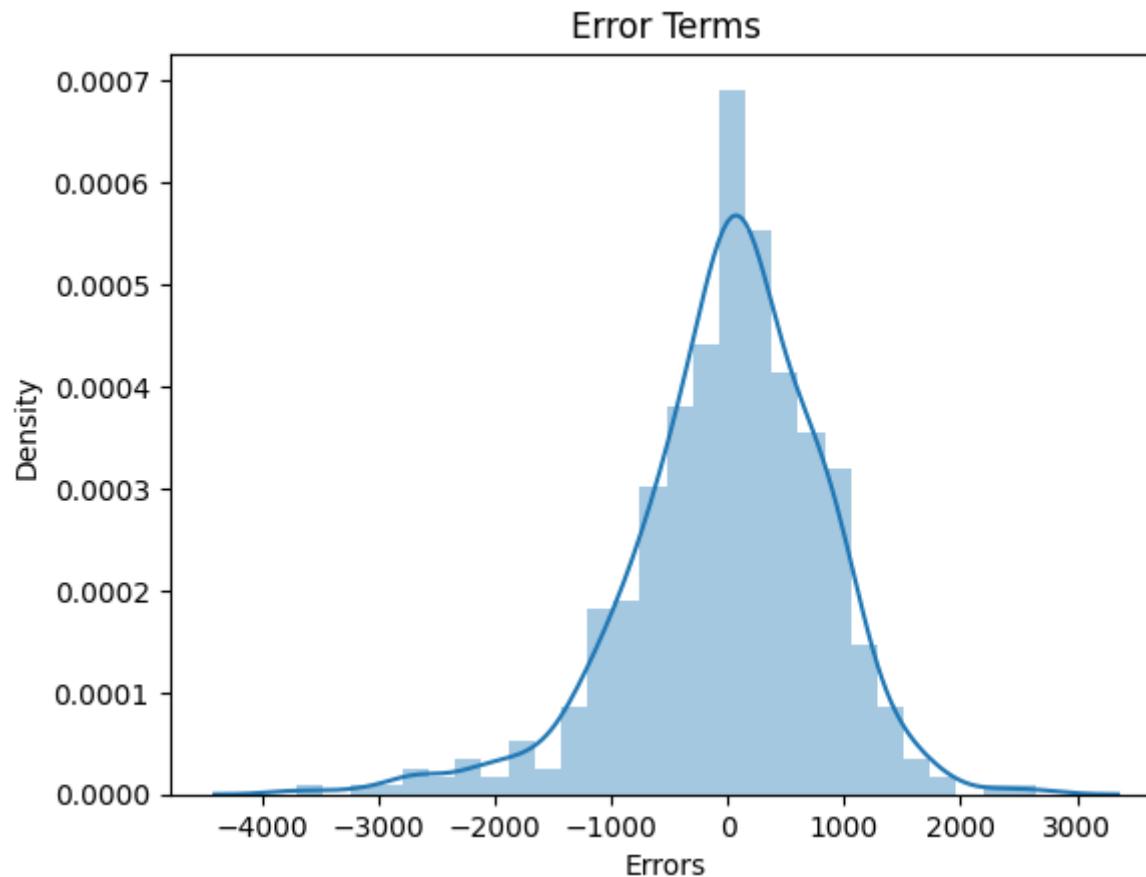
1500.5024389131136 [ 1989.45381428 -744.17786846  865.51084562  462.12893069
 733.22170346 -825.23454868 -204.12695705 -2289.06746494
-678.42373184 4091.51711769]
```

Residual Analysis

```
In [55]: y_train_pred = lr.predict(X_train[cols6])
```

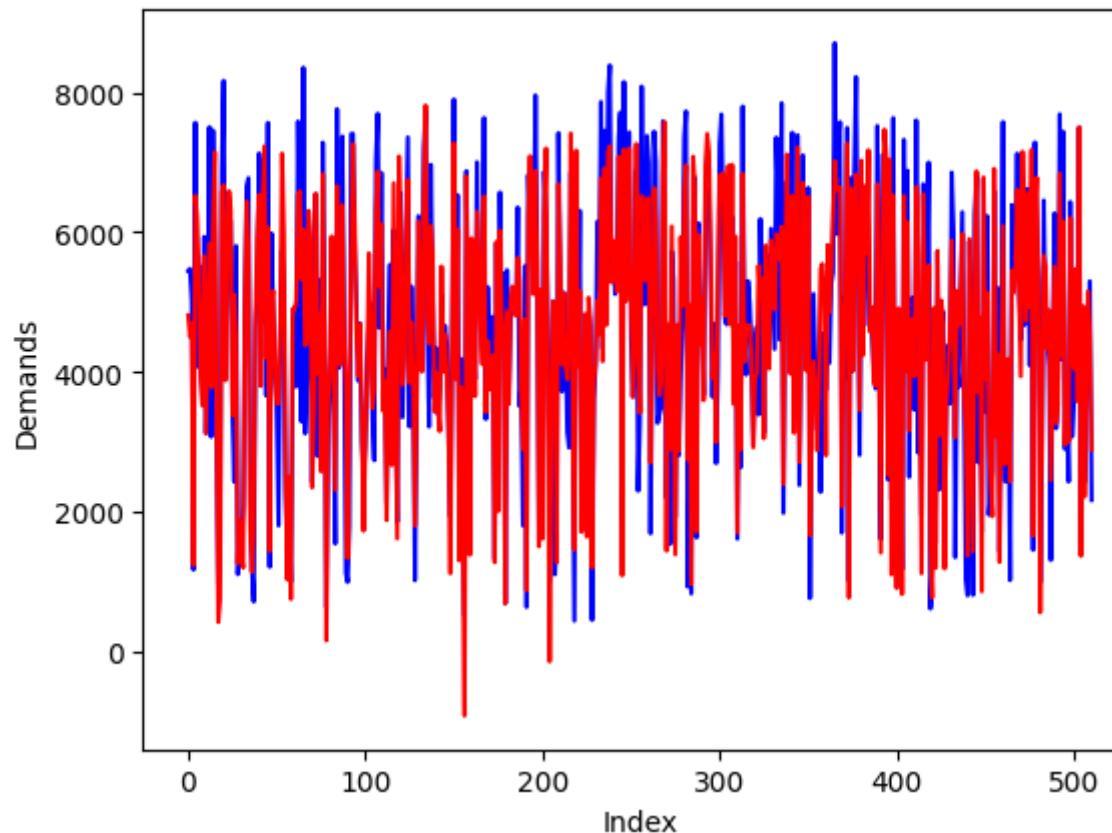
```
In [56]: #Plot a histogram of the error terms
def plot_res_dist(act, pred):
    sns.distplot(act-pred)
    plt.title('Error Terms')
    plt.xlabel('Errors')
```

```
In [57]: plot_res_dist(y_train, y_train_pred)
```



```
In [58]: # Actual vs Predicted
c = [i for i in range(0,len(X_train),1)]
plt.plot(c,y_train, color="blue")
plt.plot(c,y_train_pred, color="red")
plt.suptitle('Actual vs Predicted', fontsize = 15)
plt.xlabel('Index')
plt.ylabel('Demands')
plt.show()
```

Actual vs Predicted



Model Predictions

R-Squared value for train predictions

```
In [59]: #Print R-squared Value  
r2_score(y_train,y_train_pred)
```

Out[59]: 0.8207866077552256

Prediction of values on test dataset

```
In [60]: #Scale variables in X_test
num_vars = ['temp','atemp','hum','windspeed']

#Test data to be transformed only, no fitting
X_test[num_vars] = scaler.transform(X_test[num_vars])
```

```
In [63]: cols6 = ['yr','season_spring',
             'season_winter','season_summer','mnth_sept','holiday','workingday',
             'weathersit_bad', 'weathersit_moderate','temp']

#Predicting test data values
y_test_pred = lr.predict(X_test[cols6])
```

R-Squared value for test predictions

```
In [64]: # Find out the R squared value between test and predicted test data sets.
r2_score(y_test,y_test_pred)
```

```
Out[64]: 0.8141443198305731
```

Evaluating the model

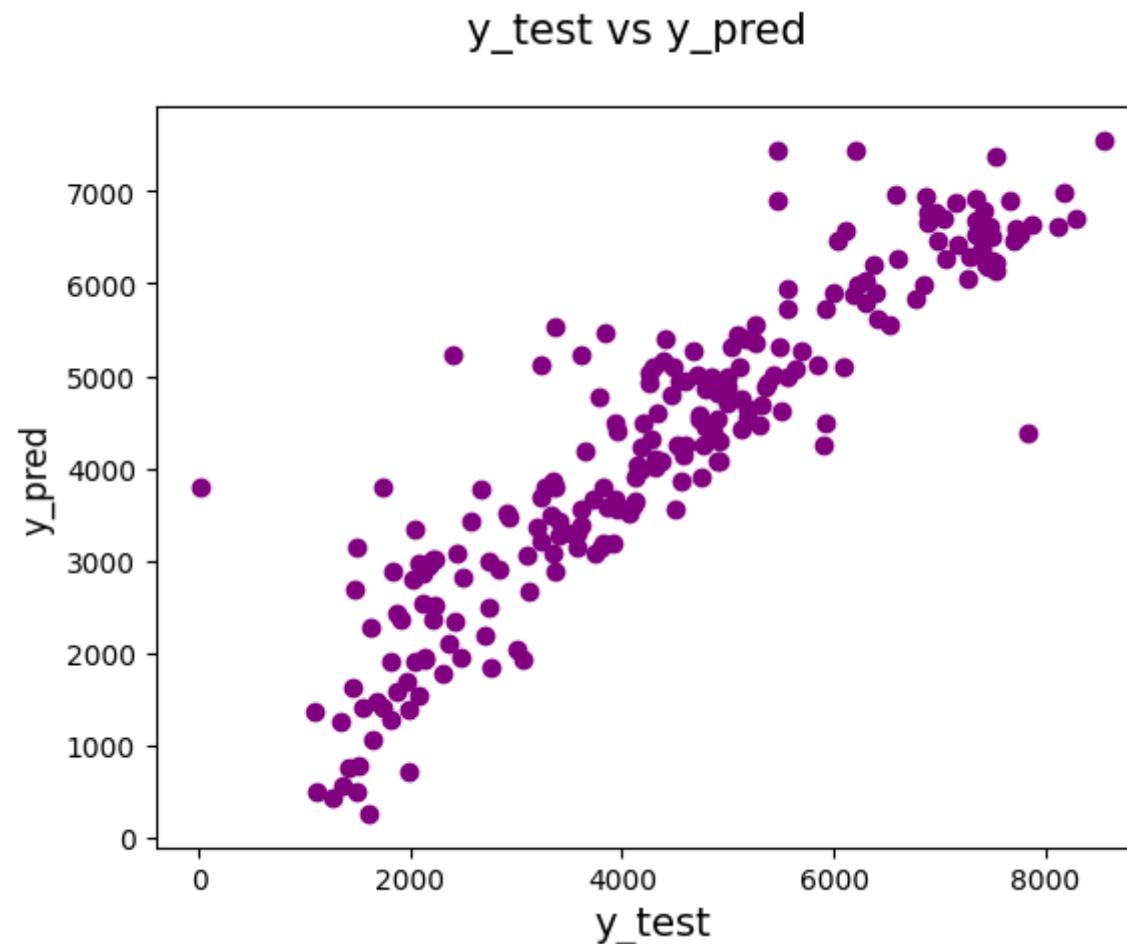
Evaluating the model based on r2_score(as mentioned in the problem statement as well)

```
In [65]: y_pred=y_test_pred
r2_score(y_test, y_pred)
```

```
Out[65]: 0.8141443198305731
```

```
In [68]: fig = plt.figure()
plt.scatter(y_test, y_pred,color='purple')
fig.suptitle('y_test vs y_pred', fontsize = 15)
plt.xlabel('y_test', fontsize = 14)
plt.ylabel('y_pred', fontsize = 12)
```

```
Out[68]: Text(0, 0.5, 'y_pred')
```



Let us rebuild the final model of manual + RFE approach using statsmodel to interpret it:

```
In [70]: cols6 = ['yr', 'season_spring',
   'season_winter', 'season_summer', 'mnth_sept', 'holiday', 'workingday',
   'weathersit_bad', 'weathersit_moderate', 'temp']

lm = build_model(cols6)
```

OLS Regression Results

Dep. Variable:	cnt	R-squared:	0.821			
Model:	OLS	Adj. R-squared:	0.817			
Method:	Least Squares	F-statistic:	229.0			
Date:	Tue, 01 Aug 2023	Prob (F-statistic):	1.70e-179			
Time:	13:43:41	Log-Likelihood:	-4153.4			
No. Observations:	511	AIC:	8329.			
Df Residuals:	500	BIC:	8375.			
Df Model:	10					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	1500.5024	262.410	5.718	0.000	984.940	2016.065
yr	1989.4538	73.778	26.965	0.000	1844.500	2134.408
season_spring	-744.1779	189.384	-3.929	0.000	-1116.264	-372.092
season_winter	865.5108	145.733	5.939	0.000	579.187	1151.834
season_summer	462.1289	122.140	3.784	0.000	222.159	702.099
mnth_sept	733.2217	138.143	5.308	0.000	461.809	1004.634
holiday	-825.2345	233.487	-3.534	0.000	-1283.970	-366.499
workingday	-204.1270	82.866	-2.463	0.014	-366.935	-41.319
weathersit_bad	-2289.0675	221.187	-10.349	0.000	-2723.638	-1854.497
weathersit_moderate	-678.4237	78.385	-8.655	0.000	-832.429	-524.419
temp	4091.5171	295.064	13.867	0.000	3511.798	4671.236
Omnibus:	63.271	Durbin-Watson:	2.151			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	118.647			
Skew:	-0.734	Prob(JB):	1.72e-26			
Kurtosis:	4.848	Cond. No.	19.0			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Conclusion

Significant variables to predict the demand for shared bikes

- yr
- season_spring
- season_winter
- season_summer
- mnth_sept
- holiday
- workingday
- weathersit_bad
- weathersit_moderate
- temp

On the Final model on train and test dataset $R^2 = 0.8$ and P-Value is close to 0.0.

In []:

In []:

In []:

In []: